
Table of Contents

No table of contents entries found.

1. Introduce yourself

Hi, thanks for having me. I am Sudheer Kumar M, a Lead Data Engineer with over 12 years of IT experience, specializing in cloud-native data engineering, automation, and analytics solutions.

currently, I am working at Kahramaa - Qatar General Electricity and Water Corporation (Nov 2023 – Feb 2025) as a Data & Business Analyst focusing on Azure DevOps data pipelines. Here, I develop scalable ETL pipelines using Azure Data Factory and Databricks, integrating data from on-premises and cloud sources, including IoT devices. I have implemented Lakehouse architectures in ADLS Gen2 and enabled real-time analytics with Databricks Structured Streaming. I also ensure data governance and security compliance, building self-service Power BI dashboards for business users.

Prior to that, I served as a Lead Data Engineer at Kyndryl India Pvt Ltd (Mar 2022 – Nov 2023) on a critical automation project for National Australian Bank. I designed a data Lakehouse architecture with Azure Data Factory, Delta Lake, and Databricks, integrating real-time telemetry and ITSM data from ServiceNow. I engineered optimization workflows using PySpark and OR-Tools to maximize change window utilization and automated over 40% of low-risk change requests with high success rates.

Before Kyndryl, I was a Senior Data & Automation Engineer at Empower Retirement (Feb 2021 – Mar 2022), where I modernized legacy data workflows and built scalable ingestion pipelines using Databricks Auto Loader, Delta Lake, and Azure Data Lake Storage. I developed complex real-time transformation and analytics pipelines and introduced data quality checks and metadata governance with Unity Catalog.

From 2016 to 2021, I worked as a Senior Software Engineer at IBM India Pvt Ltd, delivering ETL pipelines, real-time analytics, and infrastructure automation for multiple enterprise clients. I implemented containerized microservices, integrated monitoring tools, and contributed to infrastructure as code initiatives to improve scalability, availability, and operational efficiency.

My career started at SLK Software Services as a Software Developer (2014–2016), where I worked extensively with Oracle PL/SQL, Oracle Warehouse Builder, and data integration tools to support banking applications.

Overall, I bring strong expertise in Azure data platforms, big data processing, automation, and data governance, combined with a collaborative approach to build robust, scalable data solutions aligned with business needs.

2. How does your typical day looks like?

a mix of building pipelines, monitoring data Pipelines health, and collaborating with business and IT teams.

We follow Agile, so the day starts with a quick stand-up where we sync on sprint progress and any blockers. After that, one of the first things I do is check the health of our data pipelines in Azure Data Factory — looking for failed runs, slowdowns, or any anomalies. Since we bring in data from multiple sources — Oracle ERPs, IoT sensors, streaming meters, cloud APIs — making sure data keeps landing reliably in ADLS Gen2 is critical.

I spend a good chunk of my day in Databricks, writing PySpark and Spark SQL to clean and transform data. This ranges from simple validations to complex aggregations and joins. We follow the Lakehouse medallion architecture, so I'm moving data through Bronze, Silver, and Gold layers — each with its own level of cleansing and business logic. We use Auto Loader for ingestion, which handles schema evolution automatically, so when new data formats come in, we don't have to manually intervene.

On the optimization side, I regularly run OPTIMIZE and VACUUM on Delta Lake tables to keep query performance high and storage costs down. On the serving side, I keep an eye on Synapse SQL pools to ensure Power BI reports are running fast, and I verify access controls so users only see what they're authorized to — that ties into the RBAC and governance standards we enforce in ADLS.

Security is a daily responsibility too — managing credentials through Key Vault, ensuring Private Endpoints are configured correctly, and maintaining audit trails. I also document data lineage and transformation logic as I go, because we're building this platform for long-term maintainability.

And then there's always the collaboration piece — helping business analysts translate their reporting needs into data models and ETL workflows, or coordinating with IT on infrastructure changes.

So some days are code-heavy — building new pipelines or onboarding a new data source. Other days are more about troubleshooting, optimizing, and making sure everything runs smoothly behind the scenes. That mix is what keeps it engaging."

3. Can you brief me about your project

The Business Problem:

Kahramaa is Qatar's national electricity and water corporation — the sole provider for the entire country. The project is called 'Data Transformation & Self-Service Digitization.' When I joined, most data lived in siloed Oracle and MySQL databases on-premises. Business teams had no self-service analytics capability — every report required manual effort, queries were slow, and there was no single source of truth. Additionally, Kahramaa was deploying IoT devices across their electricity and water networks — smart meters, sensors — and that telemetry data had nowhere to go for real-time analysis.

The main goal here is pretty simple: we want to make sure that all the different departments can easily and securely access a centralized Data Warehouse built on Azure Synapse Analytics. This warehouse brings together data from lots of places, so each department can find exactly what they need without any hassle.

Now, the data comes from all sorts of sources—think Oracle ERP systems, MySQL databases, IoT devices, streaming data, and real-time sensors measuring water and electricity meters. To handle all this streaming and batch data smoothly, we use Azure Event Hubs and Azure IoT Hubs for the live data, and Azure Data Factory (or ADF for short) to manage pipelines that pull in data from both on-premises systems and the cloud.

So, the pipeline has five distinct stages, and each has a clear purpose.

First, the **Landing Zone** — this is a transient staging area in ADLS Gen2 where data lands exactly as it comes from the source. CSV exports from Oracle ERP, JSON payloads from IoT smart meters, XML from APIs. No transformations here — it's purely a safe landing pad. If ingestion fails midway, we have the raw files to retry from. The folder structure is organized by source and date — so you'd see paths like /landing/oracle_erp/2025-01-15/.

From there, Databricks Auto Loader picks up the files and writes them **into Bronze** as Delta Lake tables. This is our first quality gate — we enforce schema validation, reject malformed records into a dead-letter table, and add metadata columns like ingestion timestamp, source file name, and batch ID. Bronze is strictly append-only — we never update or delete here, because it's our audit trail. For IoT streaming data, Structured Streaming reads directly from Event Hubs with watermarking for late arrivals and appends to Bronze Delta tables.


Silver is where the real transformation happens. I deduplicate using ROW_NUMBER with a window function, handle nulls and type casting, standardize code values, and join datasets together — for example, joining IoT meter readings with device metadata to enrich the data. Nested JSON from IoT payloads gets flattened into columnar format. Business rules are validated here — things like consumption readings must be positive, reading dates must be after installation dates. Any records that fail validation go to a quarantine table for review rather than blocking the pipeline. Updates come through as Delta Lake MERGE operations — upsert pattern.

Gold is the business-ready layer — dimensional models with fact and dimension tables. Pre-computed aggregations like monthly consumption by zone, billing summaries, meter health KPIs. This layer is optimized for query performance with OPTIMIZE, Z-ORDER on frequently filtered columns, and VACUUM to manage storage. This is what Power BI connects to — business users only interact with Gold."

Once the data is all processed, it feeds into our Azure Synapse Analytics data warehouse. We structure this data using a model called a star schema, which makes running analytics faster and more scalable. We protect the data with role-based access control, tied into Azure Active Directory, so users only see what they’re supposed to.

Finally, all the departments get to see the results through Power BI reports. These reports refresh frequently, almost in real-time, and respect all the security rules. That means business users can explore and analyse data on their own without waiting for IT to generate reports.

And hey, we didn’t forget about keeping an eye on everything behind the scenes. We use Azure Monitor and ADF monitoring tools to catch any issues early. Azure Purview helps us keep track of where the data came from and ensure we’re following all compliance and governance rules. Plus, we’ve tuned the system using best practices so data queries and reports are fast, even as things scale up. Of course, we have backups and disaster recovery plans in place—we want to be sure things keep running no matter what.



Layer-by-layer detail — what to say in the interview

Landing zone Temporary staging area in ADLS Gen2

Data lands here exactly as-is from the source — no transformations, no schema changes. CSV from Oracle, JSON from IoT sensors, XML from APIs, Parquet from cloud sources.

This is a transient zone — files are picked up by downstream pipelines and moved to Bronze. If ingestion fails midway, the landing zone acts as a retry buffer.

Organized by source system and date: `/landing/oracle_erp/2025-01-15/`, `/landing/iot_meters/2025-01-15/`

No Delta Lake here — just raw files. This is where ADF Copy Activities drop data before any processing.

Tools: ADF Copy Activity, Event Hubs (for streaming), ADLS Gen2 containers

Bronze (raw) Quality gate 1

Landing zone files are read by Databricks Auto Loader and converted into Delta Lake tables — this is where we get ACID transactions, schema enforcement, and time travel.

Schema validation happens here: enforce correct column names, data types, reject malformed records to a dead-letter table.

Data is append-only — we never update or delete in Bronze. Every record gets metadata columns:
`_ingestion_timestamp`, `_source_file`, `_batch_id`

Partitioned by ingestion date for efficient downstream reads.

For IoT streaming data: Structured Streaming reads from Event Hubs, applies watermarking for late-arriving data, and appends to Bronze Delta tables.

Tools: Databricks Auto Loader, Structured Streaming, Delta Lake, Event Hubs

Silver (cleansed) Quality gate 2

Deduplication: `ROW_NUMBER()` partitioned by business key, ordered by timestamp DESC, keep `rn=1`.

Null handling and data type casting: parse date strings, cast numeric fields, fill defaults where business rules allow.

Standardization: normalize code values (e.g., region codes, meter types), trim whitespace, uppercase where needed.

Joins and enrichment: join IoT readings with device metadata, join billing records with customer master data.

Flattening nested structures: IoT JSON payloads get exploded and flattened into columnar format.

Business rule validation: `fare_amount > 0`, `reading_date > installation_date`, valid `meter_type` codes.

Failed records routed to a quarantine/dead-letter table for review — don't block the pipeline.

MERGE (upsert) pattern: new records inserted, changed records updated using Delta Lake MERGE INTO.

Tools: PySpark, Spark SQL, Delta Lake MERGE, window functions (`ROW_NUMBER`, `LAG`, `LEAD`)

Gold (curated) Quality gate 3

Dimensional modeling: build fact and dimension tables — `fact_consumption`, `fact_billing`, `dim_customer`, `dim_meter`, `dim_zone`.

Pre-computed aggregations: daily/monthly consumption by zone, peak vs off-peak usage, billing summaries.

KPI tables: revenue per zone, meter health scores, water loss percentages, SLA compliance metrics.

Reconciliation checks: row counts match Silver, aggregate totals tie out, referential integrity (all `customer_ids` in `dim_customer` exist).

Optimized for query performance: `OPTIMIZE` + `Z-ORDER` on high-query columns, partitioned by date/region.

`VACUUM` to clean up old file versions and reduce storage costs (retain 7 days for time travel).

This is what Power BI and Synapse connect to — business users only see Gold layer.

Tools: PySpark, Spark SQL, Delta Lake `OPTIMIZE`/`VACUUM`/`Z-ORDER`, Databricks SQL endpoints

4. Project metrics and full-length inventory details

"Let me get specific on the architecture please refer the document.



Kahramaa Project -
Architectural Referenc



metrics.md



Kahramaa - Table
Count per Layer (Sour

🔧 Pipeline inventory

"First off, we've got about **27 ADF pipelines** in total, but they're grouped by what they do.

- **Ingestion pipelines** – around 12 of them. These handle batch stuff like pulling from Oracle ERP or MySQL CRM, either hourly or daily.
- Then we have **2 pipelines just for streaming** – they basically kick off the Databricks Structured Streaming jobs that listen to Event Hubs.
- **Orchestration pipelines** – about 6. These call Databricks notebooks to move data from Bronze → Silver and Silver → Gold.
- **Post-processing** – 4 pipelines. Things like running OPTIMIZE and VACUUM on Delta tables, or refreshing Power BI datasets.
- And finally, **error handling & monitoring** – 3 pipelines. One for reprocessing quarantined data, another for sending alerts when things fail.

Scheduling? Streaming is continuous. Batch pipelines run on cron – hourly at five past the hour, daily at 2 AM, that sort of thing. Silver transforms are event-triggered right after bronze finishes."

📊 Data volumes

"For volumes, we're talking about a national utility – about 2 million smart meters.

- **Daily raw ingestion** is around 185 GB. Most of that – about 150 GB – comes from streaming IoT data, which is roughly 192 million events per day. The rest is batch from databases and files. (880
- **Monthly raw** comes to about 5.5 TB. After we compact into Delta, Bronze is around 4.5 TB, Silver drops to 2.1 TB (because we clean and dedupe), and Gold is only about 0.9 TB per month.
- **Yearly** – raw hits 65 TB, but Gold ends up around 11 TB. That's because we aggregate 15-minute readings up to daily grains.

One fact table alone – FactElectricity Usage at daily grain – grows by about 730 million rows a year. So yes, it's big, but manageable with Synapse."

⚡ Batch vs. streaming architecture

"We're doing a **hybrid** approach – because Kahramaa needs both real-time visibility and deep historical analysis.

- **Streaming path** – IoT meters and sensors go into Event Hubs. Databricks Structured Streaming picks that up, applies watermarking for late arrivals, and appends straight to Bronze Delta tables. Latency is about 1-2 minutes. That feeds real-time dashboards and alerts – like 'this substation is about to overload'.
- **Batch path** – everything else: Oracle ERP, MySQL, flat files. ADF copies them into Landing, then Databricks Auto Loader moves them to Bronze in batch mode. From there, we do hourly or daily batch ETL to Silver and Gold. That's for billing, customer 360, compliance reports – where you don't need sub-second freshness.

Why both? Streaming gives operational speed; batch gives cost-effective heavy lifting. Best of both worlds."

🔄 Incremental strategy

"We handle increments very deliberately, layer by layer.

- **Bronze** is simple: append-only. Every new file or micro-batch just gets added. We never update or delete – it's our audit trail.
- **Silver** uses Delta MERGE – upsert pattern. For fact-like data (meter readings), we merge on the natural key plus timestamp. For dimensions like customer, it's a slowly changing dimension Type 2 – we insert a new version and flag the old one as inactive.
- **Gold** – we incrementally insert new fact rows, and dimensions are fully refreshed only when we detect changes (usually nightly).

For batch sources, we use a watermark column like last_modified_date. ADF reads only what changed since the last run. For streaming, Structured Streaming checkpoints are stored in a logs container – that gives us exactly-once semantics even if the job restarts."

🚨 Error handling

"We built error handling in three layers – so one bad record never stops the whole pipeline.

- **Layer 1 – schema validation** at Bronze. If a record is malformed (missing meter ID, negative consumption), it goes to a dead-letter table – bronze.quarantine_iot, for example. The pipeline keeps running.
- **Layer 2 – business rule validation** at Silver. Things like 'reading date must be after installation date' or 'tariff code must exist'. Failed rows move to silver.quarantine instead of failing the job.
- **Layer 3 – pipeline failures** (like a database being down). ADF retries 3 times with a 5-minute delay. If it still fails, we send an alert via Azure Monitor and page the on-call engineer.

We also have a manual reprocess pipeline – pl_quarantine_reprocess – that lets a data engineer fix the bad records and re-insert them. And everything is idempotent, so replaying a batch doesn't double count."

⚙️ Cluster configuration (Databricks)

"We've got four main clusters, each sized for its job.

- **Bronze streaming cluster** – Standard_D4ds_v5 (4 vCPU, 16 GB RAM). Auto-scales from 2 to 6 nodes. Runtime 13.3 LTS. It runs continuously.
- **Silver transform cluster** – This is memory-optimised: Standard_E8ds_v4 (8 vCPU, 64 GB). Auto-scales 4 to 20 nodes. It runs as a job cluster – spins up, does the heavy dedup and enrichment, then shuts down.
- **Gold aggregate cluster** – Standard_D8ds_v5 (8 vCPU, 32 GB), 2 to 12 nodes. Also a job cluster. We tuned `spark.sql.shuffle.partitions=400` here because the fact tables are huge.
- **Dev interactive cluster** – for data engineers. Same instance type as bronze but only 1-4 nodes, and it idles after 20 minutes.

All production clusters have autoscaling enabled, and we use Delta optimisations like `optimizeWrite` and `adaptive query execution`. Libraries include the Event Hubs connector and ABFS driver.”

“So that’s the rundown – from pipeline counts and data volumes, to how we mix batch and streaming, our incremental logic, error handling, and exactly what clusters run where.

Azure Data Factory (ADF)

5. What is ADF?

“ADF stands for Azure Data Factory. Think of it as a PaaS (platform-as-a-service) managed service in the cloud. It’s serverless and scales out automatically – so you only pay for what you use, not for idle capacity.

Its main job is threefold:

- **Data integration** – copy data from anywhere (on-prem Oracle, MySQL, APIs, flat files) to Azure.
- **Data transformation** – you can run transforms inside ADF using mapping data flows, or just orchestrate external transforms (like Databricks notebooks).
- **Orchestration** – this is the big one. ADF ties everything together into a pipeline: first copy, then run a Databricks job, then refresh Power BI, and handle failures with retries and alerts.

Other key benefits for a project like Kahramaa:

- **Integrates with external services** – Event Hubs, Databricks, Synapse, Purview, you name it.
- **Version control & deployment** – you can hook ADF to Git (Azure DevOps or GitHub) and promote pipelines across dev, test, prod.

So, in the Kahramaa architecture, ADF is the conductor of the entire orchestra – it doesn’t do all the heavy compute (that’s Databricks), but it tells everyone when to start, what to do, and what to do if something goes wrong.”

6. What are the building blocks of ADF?

“So, Azure Data Factory has several core building blocks that work together like LEGO pieces. Let me walk you through them from the perspective of our Kahramaa pipelines.”

1. Linked Service

“This is the **connection string** to an external resource – where your data comes from or goes to.

Example: a linked service to on-prem Oracle ERP (using self-hosted IR), another to ADLS Gen2, another to Azure Databricks. Without linked services, ADF can’t talk to anything outside.”

2. Dataset

“A dataset is **what** specific data you’re pointing to inside a linked service. It’s like a pointer plus a schema.

Example: `ds_oracle_ar_invoices` points to the `AR_INVOICES` table in Oracle. Or `ds_adls_landing` points to `/landing/oracle_erp/` folder in ADLS.”

3. Activity

“An activity is a **single step** in a pipeline. There are three main types:

- **Data movement** – Copy Activity (ingests from source to sink).
- **Data transformation** – Databricks Notebook activity, Stored Procedure activity.
- **Control flow** – If Condition, ForEach, Until, Wait, etc.

In Kahramaa, we use Copy Activity to pull from MySQL, then a Databricks Notebook activity to run the Silver transformation.”

4. Pipeline

“A pipeline is a **logical grouping of activities** that run as a single unit. It’s the actual workflow.

Example: `pl_ingest_oracle_erp` contains Copy Activity → then a Stored Procedure activity to log success. Pipelines can be triggered or scheduled.”

5. Trigger

“Triggers define **when** a pipeline starts.

- **Schedule trigger** – runs at 2 AM daily.
- **Tumbling window trigger** – runs every hour without overlap.

- **Event-based trigger** – when a new file lands in ADLS.
We use schedule triggers for batch ingestion, and event triggers for some real-time paths.”

6. Integration Runtime (IR)

“IR is the **compute engine** that runs activities. There are three flavours:

- **Azure IR** – fully managed in cloud, for copying between cloud services.
- **Self-hosted IR** – installed on-prem or in a VM, used to access on-prem Oracle/MySQL (critical for Kahramaa).
- **Azure SSIS IR** – for lifting SSIS packages to cloud. Not used here.”

7. Data Flow (Mapping Data Flow)

“These are **visual, code-free transformations** that run on Spark behind the scenes. You can do joins, aggregations, row filtering without Databricks. In Kahramaa, we didn’t use them much because Databricks gave us more flexibility, but they’re an option.”

8. Control Flow & Variables

“Pipelines can have **parameters** (passed in at runtime) and **variables** (set inside pipeline). You can also add **dependencies** between activities – like “this activity runs only if the previous one succeeded”.”

“So to sum up: Linked Services + Datasets = where and what. Activities + Pipeline = what you do. Triggers = when. Integration Runtime = how. And control flow = logic. That’s ADF in a nutshell.”

7. What are different activities?

1 Data Movement Activities (move data from source to sink)

- **Copy Activity** – the classic. Moves data anywhere (on-prem to cloud, cloud to cloud). *Kahramaa: Oracle ERP → ADLS Landing.*
- **Lookup Activity** – reads metadata or a small dataset (but doesn't move). *Kahramaa: get last watermark value from control table.*

2 Data Transformation Activities (change data shape/content)

- **Databricks Notebook Activity** – runs a Databricks notebook. *Kahramaa: nb_silver_deduplicate.*
- **Databricks Jar Activity** – runs a custom JAR on Databricks.
- **Databricks Python Activity** – runs a Python script on Databricks.
- **Stored Procedure Activity** – executes a stored procedure in Azure SQL/Synapse. *Kahramaa: usp_refresh_powerbi_dataset.*
- **Data Flow Activity** – runs a mapping data flow (low-code Spark transform).
- **HDInsight Activities** (Hive, Pig, MapReduce, Streaming, Spark) – legacy, rarely used now.

3 Control Flow Activities (orchestrate logic, no data movement/transform)

- **If Condition Activity** – if-else branching. *Kahramaa: if last day of month, run monthly aggregation.*
- **ForEach Activity** – loop over a collection. *Kahramaa: loop through 50 Oracle tables for ingestion.*
- **Until Activity** – repeat until condition is true. *Example: wait for file to arrive.*
- **Execute Pipeline Activity** – call another pipeline (nested). *Kahramaa: master pipeline calls child pipelines.*
- **Wait Activity** – pause for a fixed time. *Example: wait 5 minutes for cluster start.*
- **Fail Activity** – manually fail the pipeline with custom error message. *Kahramaa: if quarantine rows exceed threshold, fail and alert.*
- **Set Variable Activity** – set value of a pipeline variable. *Kahramaa: store row count or timestamp.*
- **Append Variable Activity** – add value to an array variable.
- **Filter Activity** – filter an array (not data, just in-memory). *Example: filter list of files to only those with ".csv".*
- **Join Activity** – join two arrays (in-memory). *Example: combine file list with metadata list.*

4 Metadata & Utility Activities (often overlooked but powerful)

- **GetMetadata Activity** – retrieve metadata about a dataset (file names, size, last modified, child items, etc.). *Kahramaa: scan Landing folder for new files before processing.*
- **Delete Activity** – delete files or rows from a source. *Kahramaa: clean up Landing zone after successful Bronze ingestion.*
- **Validation Activity** – wait until a dataset exists (or doesn't exist) before continuing. *Example: wait for an external system to drop a "ready" file.*
- **Until Activity** (already listed above) – often used with metadata conditions.

5 External & Extensibility Activities (custom code / APIs)

- **Azure Function Activity** – call a serverless Azure Function. *Potential Kahramaa use: call a custom function to enrich meter data with weather API.*
- **Web Activity** – call any REST API (GET, POST, PUT, etc.). *Kahramaa: after Gold load, notify downstream system via API.*
- **Webhook Activity** – similar to Web Activity but waits for a callback. *Example: wait for an external long-running job to complete.*

- **Custom Activity** – run a custom executable on an Azure Batch node (legacy).

Quick Summary Table (by category)

Category	Activities
Data Movement	Copy, Lookup
Data Transformation	Databricks (Notebook/Jar/Python), Stored Procedure, Data Flow, HDInsight (various)
Control Flow	If Condition, ForEach, Until, Execute Pipeline, Wait, Fail, Set Variable, Append Variable, Filter, Join
Metadata & Utility	GetMetadata, Delete, Validation
External & Extensibility	Azure Function, Web, Webhook, Custom

“So there you go – that’s the complete inventory. The ones we use most in Kahramaa are Copy, Lookup, Databricks Notebook, If Condition, ForEach, Execute Pipeline, Set Variable, GetMetadata, Delete, and Web. But it’s good to know the others exist for when you need them.”

8. What are the Activity dependencies?

Dependency Type	Icon Color	When does Activity B run?	Real-world example (Kahramaa)
Success	Green	Runs only if Activity A succeeds	A = Copy Oracle table to Landing. B = Databricks notebook to transform. Only run transform if copy succeeded.
Failure	Red	Runs only if Activity A fails	A = Copy activity. B = Send alert email (Web activity). If copy fails, notify the team.
Completion	Blue	Runs always – whether A succeeds, fails, or is skipped	A = Silver transformation. B = Log run status to a control table. Always log, regardless of outcome.
Skipped	Grey	Runs only if Activity A is skipped (not executed)	A = Copy file (inside an If Condition false branch). B = Log "file not found". Only log when copy is skipped.

Quick memory rule

- **Success** = happy path only.
- **Failure** = error handling only.
- **Completion** = no matter what (fire and forget).
- **Skipped** = specifically when something was bypassed.

In the ADF pipeline canvas, you set these by hovering over the green checkmark (success) on an activity and choosing the desired dependency type from the dropdown.

9. What is Az key vault? What are its uses?

What is Azure Key Vault?

Think of it as a highly secure, certified vault in the cloud. Instead of storing secrets in configuration files or source code, you place them in this vault. It integrates seamlessly with other Azure services and supports **three main types of secrets**:

- **Secrets:** Any sensitive piece of information, such as database connection strings, API keys, or passwords.
- **Keys:** Cryptographic keys used for data encryption and decryption.
- **Certificates:** For authentication, like TLS/SSL certificates, managed and renewed by the vault.

A key feature is that it can be backed by **Hardware Security Modules (HSMs)** for the highest level of protection, ensuring that even Microsoft can't view your encryption keys.

How We Use It in the Kahramaa Project

Here’s how Azure Key Vault is directly integrated into the architecture we discussed, focusing on ADF and Azure Purview:

- **Secure Connections from ADF Pipelines:** Instead of hardcoding a username and password, our pl_ingest_oracle_erp pipeline fetches them directly from Key Vault at runtime. ADF uses its **Managed Identity**, a secure, auto-managed service principal, to request these secrets, so no credentials are ever exposed in the code or logs.
- **Centralized Governance & Auditing with Purview:** By centralizing secrets in Key Vault, we can enforce strict access policies on who (or which service) can access what. This integration allows **Azure Purview** to automatically scan and map the connections between ADF pipelines and Key Vault secrets, giving us a complete, auditable lineage of data and access for compliance.

Why We Use It: Key Benefits

Using Key Vault brings several crucial advantages to our project:

- **Centralized Security:** It provides a single, secure place to store and manage all secrets, drastically reducing the risk of accidental exposure.

- **Eliminates Hard-Coded Credentials:** By removing secrets from source code and config files, we prevent them from being exposed in logs or version control systems.
- **Enhanced Access Control:** It allows for fine-grained permissions, ensuring services like ADF can only access the secrets they specifically need, not everything in the vault.
- **Simplified Management & Rotation:** Secrets can be updated or rotated in one central location. Services using them get the new values automatically, without needing to be redeployed.

This makes it a foundational piece of our security and governance framework for the Kahramaa project.

10. What is Schema inference, schema binding, schema drift?

Concept	Definition	Kahramaa Example
Schema Inference	Engine guesses schema from data.	Auto Loader reads first few IoT JSON files to infer meter readings schema.
Schema Binding	You enforce a fixed, predefined schema.	Bronze validation for Oracle ERP – rejects files with extra/missing columns.
Schema Drift	Data structure changes over time.	IoT firmware adds a new field; Delta Lake evolves table schema automatically.

“So in short: inference gives you flexibility, binding gives you control, and drift is the reality you have to manage. In Kahramaa, we use inference at the start of ingestion, binding at critical quality gates, and schema evolution (carefully) to handle drift.”

11. what is azure monitor? what are the features, benefits?

In the context of the Kahramaa project, which relies on a complex pipeline of ADF, Databricks, and Synapse, **Azure Monitor** is our single source of truth for understanding how that entire system is performing.

In short, it's Microsoft's comprehensive, unified observability service for collecting, analyzing, and acting on telemetry from all your cloud and on-premises resources.

What are its key features?

Azure Monitor brings together several powerful tools into one unified platform:

-  **Metrics:** Collects numeric time-series data (like CPU%, pipeline run duration, or Databricks job latency) for near real-time performance monitoring.
-  **Logs:** Collects detailed event and diagnostic data (like ADF failure logs, error traces) for deep analysis. The primary tool for this is **Log Analytics**, which allows you to run powerful queries across all your log data.
-  **Application Insights:** An APM (Application Performance Monitoring) tool that monitors the availability, performance, and usage of live web applications.
-  **Alerts:** Proactively notifies you when it detects issues based on metrics or logs. It can send alerts via email, SMS, webhooks, or even trigger automated actions.
-  **Visualizations:** Allows you to build rich dashboards and use native Grafana integration to visualize all your telemetry data in a single pane of glass.
-  **Insights:** Provides curated monitoring experiences for specific Azure services, like VM Insights, Container Insights, and Network Insights, offering out-of-the-box health criteria and dashboards.

What are its key benefits?

- **Unified Observability:** It gives you a single, integrated platform for all your metrics, logs, and traces, eliminating the need to switch between multiple tools.
- **Proactive Issue Detection & Faster Troubleshooting:** Smart alerts and powerful log queries help you find and fix problems before they impact users.
- **Cost Control:** By monitoring resource usage and forecasting demand, Azure Monitor can help you prevent over-provisioning and optimize cloud spend.
- **AI-Powered Insights & Integrated Security:** Integrates with AI to provide advanced insights and can help detect real-time security anomalies.
- **Seamless Integration:** Integrates natively with other key Azure services like ADF, Databricks, Synapse, Purview, and Key Vault, making it a core part of the Azure ecosystem.

In our Kahramaa project, we use Azure Monitor to track ADF pipeline health, check Databricks job durations, analyze query performance in Synapse, and set up alerts for failures. This gives us a unified view across the entire data platform, helping us ensure everything runs smoothly and efficiently.

Scenario Based questions

12. When a Spark job becomes slow in production, how do you systematically identify the exact performance bottleneck?

When a Spark job becomes slow in production, I approach it methodically rather than guessing.

First, I go to the Spark UI in Databricks and identify which stage is consuming the most time — whether it's a shuffle-heavy stage, skewed task, or excessive spill to disk.

In my current and previous Azure Databricks implementations, I've often correlated stage delays with shuffle read/write size, skewed partitions, or executor memory pressure.

Next, I check executor metrics — CPU utilization, GC time, memory usage — using Databricks logs and Azure Monitor, since I've actively monitored pipelines this way in production at KAHRAMAA and NAB. If the issue is shuffle-related, I examine join strategies and partitioning; I've resolved such issues before by using broadcast joins, repartitioning on proper keys, and enabling AQE.

If it's I/O bound, I check file sizes and small-file problems, then apply compaction or Z-ordering — which helped me reduce query latency by over 60% in high-volume workloads. My goal is always to isolate whether the bottleneck is compute, data layout, or resource contention, and fix it based on evidence from execution plans and metrics rather than assumptions.

13. If you were asked to design a batch pipeline that processes billions of records daily, how would you architect it for scale, reliability, and performance?

From my experience building large-scale pipelines in Azure using ADF and Databricks, I would design it using a Lakehouse architecture with clear Bronze–Silver–Gold separation. Data would land first in ADLS Gen2 in a raw Bronze layer using event-based or scheduled ingestion through Azure Data Factory. Processing would happen in Azure Databricks using Spark, with careful partitioning based on business keys or date to ensure parallelism at scale.

I would enable Delta Lake for ACID guarantees, schema enforcement, and time travel, which I've used extensively to maintain reliability in financial and utility domain projects.

For performance, I would optimize file sizes, use partition pruning, and apply Z-Ordering where appropriate.

Reliability would come from idempotent writes, checkpointing, proper retry strategies in ADF, (refer Q:55) and monitoring through Azure Monitor and Databricks logs. In my previous implementations, this layered and governed design helped handle high-volume workloads while maintaining stability and reducing query latency significantly.

Synapse Analytics

14. Introduction: What is Azure Synapse Analytics?

Azure Synapse Analytics is an integrated analytics service that brings together data integration, enterprise data warehousing, and big data analytics into a single, unified experience. It combines the best of SQL technologies used in enterprise data warehousing (like a dedicated SQL pool), Spark technologies used for big data (like an Apache Spark pool), and a serverless SQL endpoint to query data in a data lake. This allows you to ingest, prepare, manage, and serve data for immediate business intelligence and machine learning needs.

Here are its key components:

- **Dedicated SQL Pool (formerly SQL DW):** A data warehousing resource with predictable performance, billed based on a chosen service level (e.g., DWU500c). It's best for large-scale, structured data and high-performance reporting.
- **Serverless SQL Pool:** An on-demand query service, billed per TB of data processed. It's perfect for exploring raw data in a data lake, without provisioning any dedicated resources.
- **Apache Spark Pool:** A data engineering and data science environment, billed based on the vCore-hours used by the running pool instances. It's for data transformation, exploration, and machine learning on large datasets.
- **Data Integration (Pipelines):** A serverless data orchestration service (similar to ADF) with a code-free UI for creating Extract, Transform, Load (ETL) workflows, billed per activity run and integration runtime hour.
- **Synapse Studio:** A unified web-based UI for data engineers, data scientists, and analysts to perform data integration, data preparation, management, and analysis tasks in one place.

End-to-End Data Flow: Kahramaa Project in Action

The architecture of the Kahramaa project, as you described, is a classic implementation of a medallion lakehouse architecture. Here's how the data flows through the system:

- **Data Ingestion:** Data is ingested from two primary sources: IoT devices (smart meters and sensors) using **Azure Event Hubs** and batch sources (Oracle, MySQL) using **Azure Data Factory (ADF)**.
- **The Landing Zone:** The raw, unprocessed data from ADF is first placed into a "Landing Zone" in Azure Data Lake Storage (ADLS) Gen2. This acts as a safe, transient staging area.
- **Bronze Layer (Raw Data):** Data is read from the Landing Zone by **Databricks Auto Loader** and written to the Bronze layer as **Delta Lake tables**. This layer is an append-only audit trail, ensuring data integrity.
- **Silver Layer (Cleansed & Enriched Data):** Here, the real transformation occurs using **Databricks**. This includes deduplication, data type casting, flattening nested JSON, and joining datasets. The Silver layer uses Delta Lake MERGE for upsert operations.
- **Gold Layer (Business-Ready Data):** This layer contains the business-ready **dimensional models (star schemas)** with fact and dimension tables. Data is optimized for query performance using OPTIMIZE, Z-ORDER, and VACUUM.

- **Serving Layer (Azure Synapse Analytics):** The curated data from the Gold layer is loaded into the **Azure Synapse Analytics data warehouse** (dedicated SQL pool) for high-performance analytics.
- **Consumption & Security:** Business users access the data through **Power BI** reports, which refresh frequently. Data is protected with **role-based access control (RBAC)**, tied into Azure Active Directory.
- **Monitoring & Governance:** **Azure Monitor** is used for proactive issue detection, while **Azure Purview** provides data lineage and governance to track data origins and ensure compliance.

⚙️ Deep Dive: Dedicated SQL Pool

The dedicated SQL pool is a Massively Parallel Processing (MPP) system. When you run a query, the **Control Node** receives it, optimizes it, and passes the parallel tasks to the **Compute Nodes**. These nodes process their portion of the data in parallel and send the results back to the Control Node to be combined into the final result set.

The performance and cost of a dedicated SQL pool are determined by Data Warehousing Units (DWUs). A higher DWU means more compute power and memory.

Table Distribution Methods

How data is distributed across compute nodes is a critical design choice.

- **Hash-Distribution:** Rows are distributed based on a hash of a chosen **distribution column**. This is best for large fact tables (like FactElectricityUsage), as it can co-locate related data (e.g., all rows for one customer on the same node) for faster joins.
- **Round-Robin Distribution:** Rows are distributed evenly, but randomly, across all nodes. It's useful for staging tables where data is just being loaded and not immediately joined, as it simplifies data loading.
- **Replicated Table:** A full copy of the table is stored on every compute node. This is ideal for small dimension tables (like DimDate), as it eliminates data movement during joins, greatly speeding up queries.

Loading Data: COPY INTO, PolyBase & CETAS

- **COPY INTO:** A high-throughput T-SQL statement for loading data into a table. It's the primary and recommended strategy for its flexibility and performance.
- **PolyBase:** A technology that uses metadata to allow dedicated SQL pools to query and import data from external sources (like ADLS) using T-SQL.
- **CETAS (Create External Table As Select):** A command to export the results of a SELECT statement to an external data source, like creating a file in ADLS.

Partitioning Tables

Partitioning divides a large table into smaller, manageable segments based on a column (e.g., reading_date). This improves query performance and maintenance (e.g., swapping out old data). A common practice is to **partition fact tables by date**.

⚡ Deep Dive: Serverless SQL Pool

The serverless SQL pool is a query service over the data in your data lake. It's perfect for data exploration, logical data warehousing, and transforming data without provisioning a dedicated pool. It supports reading Parquet, CSV, JSON, and Delta Lake files.

🔒 Security & Governance

Security is built into every layer:

- **Authentication:** **Azure Active Directory** integration for centralized identity management.
- **Authorization:** **Role-based access control (RBAC)** at the workspace level and **SQL-based access control** (GRANT/DENY) for data in SQL pools.
- **Data Protection:** Features like **Row-Level Security (RLS)**, **Column-Level Security (CLS)**, **Dynamic Data Masking (DDM)**, and **Transparent Data Encryption (TDE)** to protect data at rest.
- **Governance:** **Azure Purview** integration provides a unified data map, automated data discovery, and data lineage tracking across your entire data estate.

📊 Monitoring & Performance Tuning

- **Azure Monitor:** Provides platform and application-level metrics (e.g., CPU, DWU usage) and logs.
- **Dynamic Management Views (DMVs):** Offer a programmatic way to monitor query execution, identify bottlenecks, and understand workload activity using standard T-SQL.
- **Query Insights:** A feature within the Azure portal that provides a rich, graphical view of query performance.

📄 Conclusion: Summary and Next Steps

For the Kahramaa project, you've effectively designed a modern analytics platform. Here are some next steps to deepen your understanding:

- **Tune Performance:** Identify and optimize your top 10 slowest-running queries using DMVs and Query Insights.
- **Explore Advanced Features:** Learn about materialized views, result-set caching, and workload management for optimizing concurrency.
- **Integrate Purview:** Register your Synapse workspace in Purview to automatically scan and catalog your data assets.

- **Set Up Monitoring:** Configure diagnostic settings to send Synapse logs and metrics to Log Analytics for deeper analysis and historical reporting.

I hope this detailed breakdown, anchored in the Kahramaa project, helps you understand Azure Synapse Analytics more clearly. Let me know if you'd like to dive deeper into any of the components or topics I've outlined.

Based on the Kahramaa project’s Gold layer data volume and the intended use case (Power BI dashboards, self-service analytics, high concurrency), the clear recommendation is **Dedicated SQL Pool**.

 Gold Layer Data Size Recap (Kahramaa)

Time Frame	Gold Layer (Delta) Size
Daily	~30 GB
Monthly	~0.9 TB
Yearly	~11 TB

Plus, the star schema includes large fact tables like FactElectricityUsage with ~730M rows/year.

 Recommended: Dedicated SQL Pool

Why Dedicated SQL Pool is the right choice

Factor	Dedicated SQL Pool	Serverless SQL Pool
Performance	Predictable, high-performance MPP engine. Ideal for complex star schema queries with aggregations and joins.	Query latency is higher (seconds to minutes). Not designed for interactive dashboards.
Concurrency	Supports dozens of concurrent Power BI users with workload management (queues, memory limits).	Concurrency is limited (default ~10–20 queries). High concurrency leads to throttling.
Cost model	Fixed cost based on DWU (e.g., DWU500c). Predictable monthly billing.	Pay-per-query (per TB scanned). With 30 GB/day scanned by Power BI refreshes, cost becomes unpredictable and often higher.
Data size	Optimized for 1–100+ TB data warehouses. 11 TB/year is well within ideal range.	Best for small to medium datasets (<1 TB) or ad-hoc exploration. Scanning 11 TB repeatedly is expensive.
Power BI integration	Native, high-throughput connection with DirectQuery or import mode. Supports incremental refresh.	Limited; often slow for DirectQuery over large fact tables.
Indexing & optimisation	Clustered columnstore indexes, hash distribution, partitioning – all tuned for star schema.	No indexes; relies on file partitioning and predicate pushdown only.

 Why Not Serverless SQL Pool?

1. **Performance for dashboards** – Power BI users expect sub-second to 2-second responses. Serverless queries typically take 5–30 seconds even on well-partitioned data.
2. **Cost unpredictability** – Each Power BI visual refresh could scan hundreds of GB. With 30 GB/day Gold growth, monthly scan costs would exceed dedicated SQL pool costs.
3. **No workload management** – You cannot reserve resources for critical reports. A single heavy ad-hoc query can starve other users.
4. **Limited concurrency** – With ~50 business users, serverless would throttle frequently.

 When Would Serverless SQL Pool Be Used in This Project?

Serverless is perfect for:

- **Data exploration** – Data scientists exploring raw Bronze/Silver data.
- **Schema discovery** – Analysts checking new IoT data formats before formalising Gold tables.
- **Infrequent reporting** – Monthly executive summaries that don’t require sub-second latency.
- **Cost-sensitive ad-hoc queries** – Where you don’t want to provision a dedicated pool.

But for the **Gold serving layer** that powers daily Power BI dashboards? **Dedicated SQL pool every time.**

Final Recommendation for Kahramaa

Layer	Recommended Compute
Gold (star schema)	Dedicated SQL Pool (DWU500c, scalable to DWU1000c during peak loads)
Exploration / Ad-hoc	Serverless SQL Pool (for data scientists, auditors)
Bronze/Silver queries	Serverless or Databricks SQL – depending on frequency

Bottom line: With 11 TB/year in Gold and real-time Power BI expectations, dedicated SQL pool gives you predictable performance, cost control, and enterprise features like indexing and workload isolation. Serverless is a great companion for exploration, but not for production reporting at this scale.

- **The Dedicated SQL Pool:** This is our main serving layer. It's a provisioned, high-performance engine that acts as the ultimate "Gold" layer and the source of truth for all business reporting. Its main job is to serve data to Power BI.
- **The Serverless SQL Pool:** We use this for ad-hoc exploration, to quickly query raw files directly from our Data Lake without loading them into the dedicated pool, which is great for discovery and data profiling.

Synapse analytics components

15. Discuss the components of Synapse analytics

The Main Components of Azure Synapse Analytics

Think of Synapse as a unified analytics platform that brings together **data integration, data warehousing, big data analytics, and data exploration** into one service. Here are its building blocks:

1. Dedicated SQL Pool (formerly SQL Data Warehouse)

- **What it is:** A provisioned, MPP (Massively Parallel Processing) data warehouse engine.
- **How it works:** You pay for a fixed performance level (DWU – Data Warehousing Units). The data is distributed across compute nodes.
- **Use case in Kahramaa: Gold layer serving** – high-performance star schema queries for Power BI dashboards. Our estimated 11 TB/year Gold data fits perfectly here.
- **Key features:** Clustered columnstore indexes, hash/round-robin/replicated distributions, partitioning, workload management.

2. Serverless SQL Pool

- **What it is:** An on-demand query service over files in ADLS Gen2 (Parquet, CSV, JSON, Delta).
- **How it works:** No infrastructure to provision. You pay per TB of data scanned. Queries run on the fly.
- **Use case in Kahramaa: Exploration and ad-hoc analytics** – data scientists exploring Bronze/Silver data, schema discovery, or infrequent reporting. Not recommended for high-concurrency Power BI dashboards due to latency and cost unpredictability.

3. Apache Spark Pool

- **What it is:** A managed Spark engine for big data processing, data engineering, and machine learning.
- **How it works:** You spin up a pool of nodes (provisioned or serverless). Code runs in notebooks or jobs.
- **Use case in Kahramaa: Data transformation** – though we used Databricks for Bronze→Silver→Gold, you could migrate those notebooks to Synapse Spark pools. Also useful for complex ETL, feature engineering, or running ML models.

4. Data Integration (Pipelines)

- **What it is:** A serverless, code-free orchestration service (very similar to Azure Data Factory, built on the same engine).
- **How it works:** You create pipelines with activities (Copy, Data Flow, Notebook, etc.), triggers, and dependencies.
- **Use case in Kahramaa: Orchestration** – you could replace ADF with Synapse pipelines to move data from on-prem to Landing, trigger Spark jobs, and load Gold into dedicated SQL pool. It's essentially ADF inside Synapse Studio.

5. Synapse Studio

- **What it is:** The unified web-based UI for all of Synapse.

- **How it works:** One portal where you can build pipelines, write SQL queries, develop Spark notebooks, manage lake data, monitor runs, and set up security.
- **Use case in Kahramaa: Central operations** – data engineers build pipelines, data analysts run ad-hoc SQL, and administrators monitor performance – all in one place.

6. Data Lake (ADLS Gen2)

- **What it is:** The underlying storage for Synapse. ADLS Gen2 is a hierarchical file system built on Blob storage.
- **How it works:** All Synapse components read and write from the same lake – Bronze, Silver, Gold layers live here as Delta tables or files.
- **Use case in Kahramaa: The single source of truth** – Landing, Bronze, Silver, and Gold all stored in ADLS Gen2. Synapse’s serverless SQL can query it directly; dedicated SQL pool can load from it via COPY INTO or external tables.

7. Security & Governance Integration

- **Azure Active Directory** – authentication and RBAC at workspace and SQL levels.
- **Column/Row-level security** – fine-grained access control inside dedicated SQL pool.
- **Azure Purview** – data lineage, cataloging, and governance across Synapse and other sources.

🔗 How These Components Work Together in Kahramaa

Component	Role in the Project
ADLS Gen2	Stores Landing, Bronze, Silver, Gold (Delta)
Dedicated SQL Pool	Serves Gold star schema for Power BI
Serverless SQL Pool	Used by data scientists to explore Bronze/Silver raw data
Spark Pool	(Alternative to Databricks) for Silver transformations
Pipelines	Orchestrates batch ingestion from Oracle/MySQL to Landing
Synapse Studio	Central UI for all engineering and monitoring tasks
Azure Purview	Governs lineage and compliance

🧠 Quick Comparison Table

Component	Provisioned?	Cost Model	Best For
Dedicated SQL Pool	Yes – fixed DWU	Pay per hour (fixed)	High-performance, predictable workloads (Gold layer)
Serverless SQL Pool	No – on-demand	Pay per TB scanned	Exploration, ad-hoc, infrequent queries
Spark Pool	Optional – provisioned or serverless	Pay per vCore-hour	Big data transforms, ML, data engineering
Pipelines	Serverless	Pay per activity run	Orchestration, ETL, data movement
Synapse Studio	Free (UI)	No additional cost	Unified management and development

So, in summary: Synapse is a complete analytics platform that brings together SQL, Spark, pipelines, and a lake – all accessible from a single studio. For Kahramaa, we primarily use dedicated SQL pool for Gold serving, but the other components are there for exploration, transformation, and orchestration if we ever need to consolidate away from ADF + Databricks.

16. When to use Synapse vs your existing Databricks stack

Given your Kahramaa architecture already uses ADF + Databricks + ADLS Gen2, the honest answer is: **Synapse is most additive at the serving layer**. The Dedicated SQL pool gives you a proper MPP warehouse with ColumnStore indexes and hash distribution — something Databricks SQL is also capable of, but Synapse's SQL pool is more familiar to SQL-heavy teams and integrates natively with Power BI's gateway-free DirectQuery.

Where Synapse Spark is weaker than Databricks is Delta Lake optimization (no Photon engine, no liquid clustering, no Delta Live Tables equivalent), and Unity Catalog governance is more mature. So a common hybrid pattern is: **Databricks for transformation and governance → Gold Delta tables on ADLS → Synapse Dedicated SQL pool for BI serving**.

Spark Related topics

17. How Spark decides which join strategy to pick?

Spark's **Catalyst optimizer** is responsible for this decision. It goes through a multi-phase process before any physical execution happens.

The decision hierarchy (in order)

Spark evaluates join strategies in a strict priority order:

- 1. Broadcast Hash Join (BHJ)** — checked first The optimizer checks if either side of the join is small enough to broadcast to all executors. The threshold is controlled by spark.sql.autoBroadcastJoinThreshold (default **10 MB**). If the table's estimated size (from statistics) is below this threshold, Spark broadcasts it — no shuffle required.

This is the fastest join. In your Kahramaa context, dim_meter or dim_location would almost always qualify, while fact_meter_readings never would.
- 2. Sort Merge Join (SMJ)** — the default for large tables If neither side is small enough to broadcast, Spark falls back to SMJ. Both datasets are shuffled on the join key, sorted, and then merged. It's disk-spill tolerant, which makes it the safe default for large-to-large joins.
- 3. Shuffle Hash Join (SHJ)** — situational If one side is significantly smaller than the other but still too large to broadcast, Spark may choose SHJ — it shuffles both sides on the join key but builds a hash table from the smaller side instead of sorting. It's faster than SMJ when memory allows, but risky if the build side is larger than expected (can cause OOM). Spark is conservative about choosing this.
- 4. Broadcast Nested Loop Join (BNLJ)** — last resort Used when there's no equi-join condition (e.g., a cross join or a non-equality predicate like a.ts BETWEEN b.start AND b.end). Very expensive — $O(n^2)$ — Spark will warn in the query plan.
- 5. Cartesian Product Join** Only when you explicitly write a cross join with no join condition. Spark requires you to set spark.sql.crossJoin.enabled = true as a safety gate.

What the Catalyst optimizer actually does

The decision isn't just about table size — it's a cost-based process:

Statistics collection is the foundation. If you've run ANALYZE TABLE (or in Delta Lake, statistics are auto-collected on write), Catalyst has row counts, column-level min/max, and null counts. Without statistics, it falls back to heuristics based on file sizes.

Logical plan → Physical plan. Catalyst first builds a logical plan (what joins need to happen), then the **Cost-Based Optimizer (CBO)** evaluates multiple physical plans and picks the lowest estimated cost. CBO is enabled via spark.sql.cbo.enabled = true — it's off by default in older versions.

Partition pruning and predicate pushdown happen before join strategy selection — so by the time Spark is deciding between BHJ and SMJ, it's already filtered out partitions you don't need. This can change table size estimates significantly.

How you can influence the decision

Method	How
Hint in SQL	SELECT /*+ BROADCAST(dim_meter) */ ...
Hint in PySpark	df.hint("broadcast")
Raise broadcast threshold	spark.conf.set("spark.sql.autoBroadcastJoinThreshold", "50mb")
Force SMJ	df.hint("merge")
Force SHJ	df.hint("shuffle_hash")
Enable CBO	spark.conf.set("spark.sql.cbo.enabled", "true")
Collect stats	ANALYZE TABLE fact_meter_readings COMPUTE STATISTICS FOR ALL COLUMNS

How to read the actual decision

```
python
df.explain(mode="formatted")
```

Look for the physical operator in the output — BroadcastHashJoin, SortMergeJoin, BroadcastNestedLoopJoin. If you see SMJ where you expected BHJ, it usually means statistics are missing or stale, and Spark overestimated the table size.

Interview-ready answer (spoken version)

"Spark's Catalyst optimizer decides the join strategy based on estimated table sizes derived from statistics. It first checks if either side falls below the broadcast threshold — 10 MB by default — in which case it picks Broadcast Hash Join, which avoids a shuffle entirely. If neither side qualifies, it defaults to Sort Merge Join, which shuffles and sorts both sides on the join key. Shuffle Hash Join is chosen when one side is smaller but not small enough to broadcast, though Spark is conservative about this due to OOM risk. For non-equi joins, it falls back to Broadcast Nested Loop Join. The decision quality depends heavily on whether statistics are collected — without them, Spark uses file size heuristics which can misfire. You can override the decision using query hints like BROADCAST or MERGE, or by tuning autoBroadcastJoinThreshold. In Delta Lake specifically, statistics are written at the file level on every write, which gives Catalyst much better size estimates than plain Parquet."

That last Delta Lake point is a good differentiator in a senior interview — it shows you understand why Delta + Spark performs better than raw Parquet + Spark even with identical data.

18. can you explain Spill-to-disk despite enough memory


Why Spark spills to disk even when memory looks sufficient

The key insight the interviewer wants to hear: **"enough total memory" and "enough memory for the right thing at the right time" are two very different things.** Spark's memory is divided into regions with fixed fractions, and a spill happens when one specific region is exhausted — regardless of what's free elsewhere.

Spark's unified memory model (the root cause)

Since Spark 1.6, executor memory is partitioned as follows:

Total executor memory



Execution and Storage share the unified pool and can borrow from each other — but only Storage can evict from Execution, not the other way around. **If execution memory is under pressure, it spills — even if storage memory is sitting idle and the overall JVM heap looks healthy**

The simple mental model first

Think of each Spark task as a worker at a desk. The desk has a fixed amount of workspace (execution memory). When the work piles up beyond what fits on the desk, the worker has to temporarily stack papers on the floor (disk) to make room — that's spill. The floor is slow. The desk is fast.

The problem is never "we don't have enough office space overall" — it's always "this one worker got handed too much work for their desk size."

Scenario 1 — One worker gets all the work (data skew)

What happens: You're joining meter readings on meter_id. 80% of your data belongs to one big substation. So one task gets 80% of the rows, its desk overflows, and it spills. The other 199 tasks finish in 2 seconds. That one task runs for 20 minutes.

How you'd spot it: Spark UI shows one task taking 10× longer than all others, with a big "Spill (Disk)" number next to it.

Fix — Salting: You artificially spread the skewed key across multiple tasks by adding a random suffix to it.

python

import pyspark.sql.functions as F

Add a random salt 0-9 to the skewed side

```
fact = fact.withColumn("meter_id_salted",
    F.concat(F.col("meter_id"), F.lit("_"), (F.rand() * 10).cast("int")))
```

Explode the dim side to match all 10 salt values

from pyspark.sql.functions import array, explode

```
dim = dim.withColumn("salt", explode(array([F.lit(i) for i in range(10)])))
```

```
dim = dim.withColumn("meter_id_salted",
    F.concat(F.col("meter_id"), F.lit("_"), F.col("salt")))
```

```
# Now join on the salted key — work is spread evenly
```

```
result = fact.join(dim, "meter_id_salted")
```

Or just turn on AQE and let Spark detect and split skewed partitions automatically:

```
python
```

```
spark.conf.set("spark.sql.adaptive.skewJoin.enabled", "true")
```

Scenario 2 — Too much data per task (too few partitions)

What happens: Your default shuffle partition count is 200. But you have 400 GB of data. So each task gets 2 GB to sort/hash, but the desk only fits 500 MB. Spill is unavoidable — not because of skew, but because every single task is overloaded.

How you'd spot it: Every task (not just one) shows spill in the Spark UI. Shuffle read size is uniformly large across all tasks.

Fix — Increase shuffle partitions:

```
python
```

```
# Rule of thumb: aim for ~200MB per partition
```

```
spark.conf.set("spark.sql.shuffle.partitions", "2000")
```

```
# Or let AQE figure it out automatically
```

```
spark.conf.set("spark.sql.adaptive.coalescePartitions.enabled", "true")
```

With AQE enabled, Spark starts with a large number of partitions and automatically merges small ones — so you don't have to guess.

Scenario 3 — A broadcast variable eating into the space (memory region conflict)

What happens: You broadcast a dimension table (say `dim_location`). That broadcast copy lives in storage memory. Now a big Sort Merge Join tries to use execution memory — but storage memory has quietly eaten into the shared pool. Execution memory is squeezed, spill happens.

How you'd spot it: Broadcast size in Spark UI is large (hundreds of MB). Spill happens on joins that follow the broadcast, even though total memory looks fine.

Fix — Either reduce what you broadcast, or increase the memory fraction:

```
python
```

```
# Option 1: Don't broadcast large tables — remove the hint :::: # df.hint("broadcast") -- remove this if dim table > 100MB
```

```
# Option 2: Give execution memory a larger share of the pool :::: spark.conf.set("spark.memory.fraction", "0.7") # default is 0.6
```

Scenario 4 — Window functions buffering everything

What happens: You write something like "running total of consumption per meter". Spark has to hold all rows for that meter in memory before it can emit any result. If one meter has 10 million rows, that entire buffer has to fit on one task's desk.

```
python
```

```
# This forces Spark to buffer ALL rows per meter_id
```

```
from pyspark.sql.window import Window
```

```
import pyspark.sql.functions as F
```

```
w = Window.partitionBy("meter_id").orderBy("reading_ts")
```

```
df = df.withColumn("running_total", F.sum("consumption").over(w))
```

Fix — Pre-aggregate to reduce rows before the window:

```
python
```

```
# Instead of running a window over raw readings,
```

```
# aggregate to hourly first, then apply the window
```

```
hourly = df.groupBy("meter_id", F.hour("reading_ts")).agg(F.sum("consumption").alias("hourly_consumption"))
```

```
w = Window.partitionBy("meter_id").orderBy("hour")
```

```
result = hourly.withColumn("running_total", F.sum("hourly_consumption").over(w))
```

You've gone from 10 million rows per meter down to 8,760 (hours in a year) — the buffer is now tiny.

The one setting that helps all scenarios

python

```
spark.conf.set("spark.sql.adaptive.enabled", "true")
```

AQE is like giving Spark a live feed of what's actually happening — partition sizes, skew, spill risk — and letting it re-plan mid-execution. It handles scenarios 1 and 2 automatically. Think of it as the safety net underneath all your other tuning. In Databricks (your Kahramaa stack), it's on by default. In vanilla Spark, you need to enable it explicitly.

Quick cheat sheet

Symptom	Cause	Fix
One task takes 10× longer, big spill	Data skew	Salt the key or enable AQE skew join
All tasks spilling uniformly	Too few partitions	Increase shuffle partitions or enable AQE coalescing
Spill after a broadcast	Memory region conflict	Reduce broadcast size or raise memory.fraction
Slow window function on large partition	Buffer overflow	Pre-aggregate before the window

The interviewer isn't expecting you to memorise config values — they want to see that you can **read the symptom, trace it to a root cause, and name the right lever**. That's the whole answer.